

Machine Learning Assignment 2: Neural Networks

Deadline: Thursday 23rd November 2023

Please submit your solution in PDF format (preferably, but not necessarily, L^AT_EX—this .tex file can be found on iCorsi). Handwriting and scanned documents are not allowed. In case you need further help, please write on iCorsi or contact me at vincent.herrmann@usi.ch.

Question 1 (20 points). A two-layer neural network is defined by the following equation:

$$y_k := \mathbf{w}^{(2)} f(\mathbf{x}_k W^{(1)} + \mathbf{b}^{(1)})^T + b^{(2)}$$

The values of the input vectors \mathbf{x}_k (summarized in the matrix X), targets t_k (summarized in the vector \mathbf{t}), weights $W^{(1)}$ and $\mathbf{w}^{(2)}$ as well as biases $\mathbf{b}^{(1)}$ and $b^{(2)}$ are given below.

$$X = \begin{bmatrix} - & \mathbf{x}_1 & - \\ - & \mathbf{x}_2 & - \\ - & \mathbf{x}_2 & - \\ - & \mathbf{x}_2 & - \end{bmatrix} = [[0.6, -1.0], [0.8, -1.0], [-0.4, 0.9], [0.2, 0.0]]$$

$$\mathbf{t} = \begin{bmatrix} t_1 \\ t_2 \\ t_3 \\ t_4 \end{bmatrix} = [-0.8, -0.1, 0.9, 0.7]$$

$$W^{(1)} = [[-0.8, -0.7, 0.6], [-1.0, 0.5, -1.0]]$$

$$\mathbf{b}^{(1)} = [-0.2, -1.0, -0.7]$$

$$\mathbf{w}^{(2)} = [0.1, -1.0, 0.5]$$

$$b^{(2)} = -0.7$$

The function f is the ReLU nonlinearity. The loss of the model is

$$L := \frac{1}{2} \sum_{k=1}^4 (y_k - t_k)^2.$$

Numerically calculate the gradients of L with respect to $W^{(1)}$, $\mathbf{w}^{(2)}$, $\mathbf{b}^{(1)}$ and $b^{(2)}$, and explain your process. You may use a calculator or math software (numpy, matlab, ...), but no auto-differentiation libraries like PyTorch or Jax. Tip: If you use matrix-matrix multiplications involving X , you don't need to do multiple explicit forward and backward passes.

Question 2 (15 points). As mentioned in the lecture and on the slides, often the Softmax nonlinearity at the output neurons and the Crossentropy loss are combined into a single Softmax+Crossentropy loss function. With this loss function, the error is $E = -\sum_k t_k \log \left(\frac{\exp(y_k)}{\sum_i \exp(y_i)} \right)$, with y_k being the inputs to the softmax function (i.e. the outputs of the neural network), and t_k being the elements of the target vector. Derive and simplify the gradient $\frac{\partial E}{\partial y_i}$. How can it be simplified even more if the target is a one-hot vector (one element has the value 1, all the others are zero)?

Question 3 (15 points). If we choose a learning rate that is too high, gradient descent can diverge (even in the case of convex functions). Assume we want to use gradient descent to find the value of x that minimizes $f(x) = 5x^2 + 2$. The update via gradient descent can be written as $x_{n+1} = x_n - \eta f'(x_n)$, with η being the learning rate. What is the range of values that η can take so that gradient descent converges to the minimum from any starting point? Provide an explanation of your result (can be visual).